

We thank the reviewers for their time and thoughtful feedback. Taking their helpful comments into account, we have sought to clarify the presentation of our work. Besides clarifying our current result for multiple tasks, we have also extended our result on variance reduction for multiple tasks in response to a common criticism by R1 and R2 (L11-17).

Predicting a particular task using MTL? [R2] As R2 pointed out, we focus on a situation where we have a target task for which we only have limited labeled data and a source task. We study when training the tasks together can benefit the target task. While this setting differs from traditional MTL that studies the average performance of all tasks, it is also a common setting for MTL in practice. For example, for predicting a rare event or classifying a Xray-scan, collecting large amounts of labeled data is not possible or very expensive, but auxiliary labeled data are often easier to obtain. Traditional MTL theory that studies the average performance of all tasks does not predict whether using MTL can benefit the target task. Our work applies to this setting and takes a step towards filling the gap.

What can we say for multiple tasks? [R1, R2] We have focused on two tasks in the submission to provide insights, since this is the foundational setting. In addition, we have *extended our result on bias-variance tradeoff to multiple tasks* as well. **(1)** We can now show that *as long as the output dim. of the shared layer B is smaller than the total number of tasks, the variance of the MTL estimator for the target task is always smaller than the variance of the STL estimator but the bias is always larger*. The idea is similar to two tasks and we have included this result in the draft. **(2)** For multi-label settings where all tasks have the same features, i.e. $X_i = X$ for any i , using Thm 3.6 *all of our insights into two tasks still apply except for the covariate shift* (covariate shift doesn't apply since tasks have the same features).

Writing: [R2, R3] We have corrected the typos that R2 pointed out and clarified the issues that R3 raised. **(1)** L112-118: we use t to denote the number of tasks, hence for two tasks $t = 2$. **(2)** Validation set size L108: we only need it to be larger than the size of the hidden layer times the number of tasks, which is much smaller compared to the size of the training set. **(3)** Def. of the prediction loss L113: the expectation is over a test sample x whose label is $x^\top \beta_t$. Taking expectation over ε gives the bias-variance decomposition, following standard linear regression literature [17,18].

R1: (♦) We thank R1 for suggesting looking at qualitative predictions of Thm 3.6 for multiple tasks, which we have added in the draft. **(1)** For task similarity, the more similar tasks are, the more variance reduces ($\|v_t\|$ closer to 1), which leads to positive transfer as in Prop 3.3. **(2)** For sample ratio, the more dissimilar tasks are, the more bias increases w/ more source samples, which leads to negative transfer as in Prop 3.4. **(♦)** R1 asks how our method compares to loss reweighting. Our method is equivalent to increasing task weight until performance drops. An advantage of our method is that we only compute over a subset of samples whereas loss reweighting uses the full set. **(♦)** We thank R1 for pointing out the vague use of "similar performance" in experiments, which we replaced w/ (comparable) accuracy numbers.

(♦) R1 asks how do the bounds differ from previous theory. The closest work to ours is [15] and that work uses standard concentration bounds to show that when two tasks are similar enough, MTL ensures positive transfer. Our work uses new tools from random matrix theory, and Thm 3.2 doesn't require tasks to be similar. Using the tools we *rigorously study the phenomenon of negative transfer* including *varying sample size and covariate shift, all of which aren't possible using standard concentration bounds*. **(♦)** R1 suggests computing similarity via distance between layer parameters, which we tried (along w/ SVCCA) as a proxy for task similarity but it didn't perform as well as our result in Table 1. Our framework also studied varying sample size and covariate shift in MTL, both of which are not known in prior work.

R2: (♦) We thank R2 for bringing up the confusion of which sample size regime our theory/algorithm applies, which we have clarified in the draft. **(1)** R2 is correct that "our theory applies when the sample sizes are 10-100x of feature dim.". We think this is a reasonable regime to consider; for example, in our sentiment analysis experiment, the feature dim. of a sentence is 300 and the training set size ranges from 3k to 10k. **(2)** R2 mentions "having imbalanced sample size btw source/target task": We will clarify the writing after Thm 3.2 but our incremental training scheme does not assume that the tasks have imbalanced sample size; for example, in our sentiment analysis experiment, we have observed that our method can help even when the source task is smaller than the target. The correct way to think about "imbalance" is that it also depends on task similarity, and it can be provably small (e.g. ≤ 1 cf. Prop 3.4).

(♦) We thank R2 for pointing out the connection between our incremental training procedure and curriculum learning. We are not aware of any previous work that proposes such an idea in MTL while having a strong theoretical basis. Adding more context, there is an ongoing discussion of how much data from each task the model should be trained on (cf. Google T5 and references therein). We have focused on evaluating training efficiency as a *further validation of our theory*. It's conceivable that by combining our procedure w/ other ideas one might get better final performance of the target task. As R1 pointed out, we think this is an interesting research question for future study.

R3: We thank R3 for commenting on our work. **L108:** As stated on L108 "a validation set that's larger than $r \cdot t \leq t^2$ suffices", we will further clarify here as in L19 above ($p^{0.99}$ can be replaced w/ $p^{0.5}$). **L113:** The use of t : We have used t to denote the number of tasks as well as the index of the target task. **L117:** The sample covariance of task 1 is $X_1^\top X_1$ not Σ_1 . **L187:** Parameter of the metric: yes one can recover the entire precision-recall curve by varying it. **L220:** Value of our theory to our algorithm: For two tasks, we can show that our algorithm can provably find the optimal sample ratio. As shown in Fig 1b, the performance curve, which is a quadratic function, has a single peak and our algorithm stops at the peak. That the curve is quadratic is shown in our proof and we have added the connection to the draft.